

# Unsupervised Learning - Final Project by Marshall Folkman

June 26, 2023

## 1 Unsupervised Learning - Final Project by Marshall Folkman

### 1.0.1 Project Description:

The goal of this exercise is to predict the top ten poverty stricken countries that are in greatest need for humanitarian aid. We will do this by performing PCA and building a K-means unsupervised learning model that can be used for future datasets as the world evolves.

**Data Source Citation:** <https://www.kaggle.com/datasets/rohan0301/unsupervised-learning-on-country-data?select=Country-data.csv>

**Github Source:** <https://github.com/Vamboozzer/AI/tree/main/UnSupervisedML/CountryData-MostPoverty>

**Video Presentation:** <https://youtu.be/xVJShH-omXs>

### 1.1 Part 1: Data Cleaning, EDA, and PCA

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.ensemble import IsolationForest

url_CountryData = 'https://raw.githubusercontent.com/Vamboozzer/AI/main/
↳UnSupervisedML/CountryData-MostPoverty/Country-data.csv'
url_FeatureDefinitions = 'https://raw.githubusercontent.com/Vamboozzer/AI/main/
↳UnSupervisedML/CountryData-MostPoverty/data-dictionary.csv'

CountryData = pd.read_csv(url_CountryData)
FeatureDefinitions = pd.read_csv(url_FeatureDefinitions)
```

```
[2]: # display feature definitions
pd.set_option('display.max_colwidth', None)
feature_defs_styled = FeatureDefinitions.style.set_properties(**{'text-align': 'left'})
display(feature_defs_styled)
```

<pandas.io.formats.style.Styler at 0x7f1c6e1b1e10>

```
[3]: # Overview of dataset
print("Shape of Data: ", CountryData.shape)
display(CountryData.head(12))
```

Shape of Data: (167, 10)

	country	child_mort	exports	health	imports	income \
0	Afghanistan	90.2	10.0	7.58	44.9	1610
1	Albania	16.6	28.0	6.55	48.6	9930
2	Algeria	27.3	38.4	4.17	31.4	12900
3	Angola	119.0	62.3	2.85	42.9	5900
4	Antigua and Barbuda	10.3	45.5	6.03	58.9	19100
5	Argentina	14.5	18.9	8.10	16.0	18700
6	Armenia	18.1	20.8	4.40	45.3	6700
7	Australia	4.8	19.8	8.73	20.9	41400
8	Austria	4.3	51.3	11.00	47.8	43200
9	Azerbaijan	39.2	54.3	5.88	20.7	16000
10	Bahamas	13.8	35.0	7.89	43.7	22900
11	Bahrain	8.6	69.5	4.97	50.9	41100

	inflation	life_expec	total_fer	gdpp
0	9.440	56.2	5.82	553
1	4.490	76.3	1.65	4090
2	16.100	76.5	2.89	4460
3	22.400	60.1	6.16	3530
4	1.440	76.8	2.13	12200
5	20.900	75.8	2.37	10300
6	7.770	73.3	1.69	3220
7	1.160	82.0	1.93	51900
8	0.873	80.5	1.44	46900
9	13.800	69.1	1.92	5840
10	-0.393	73.8	1.86	28000
11	7.440	76.0	2.16	20700

### 1.1.1 EDA

```
[4]: # Summary statistics
display(CountryData.describe())
```

	child_mort	exports	health	imports	income \
count	167.000000	167.000000	167.000000	167.000000	167.000000
mean	38.270060	41.108976	6.815689	46.890215	17144.688623
std	40.328931	27.412010	2.746837	24.209589	19278.067698
min	2.600000	0.109000	1.810000	0.065900	609.000000
25%	8.250000	23.800000	4.920000	30.200000	3355.000000
50%	19.300000	35.000000	6.320000	43.300000	9960.000000
75%	62.100000	51.350000	8.600000	58.750000	22800.000000
max	208.000000	200.000000	17.900000	174.000000	125000.000000

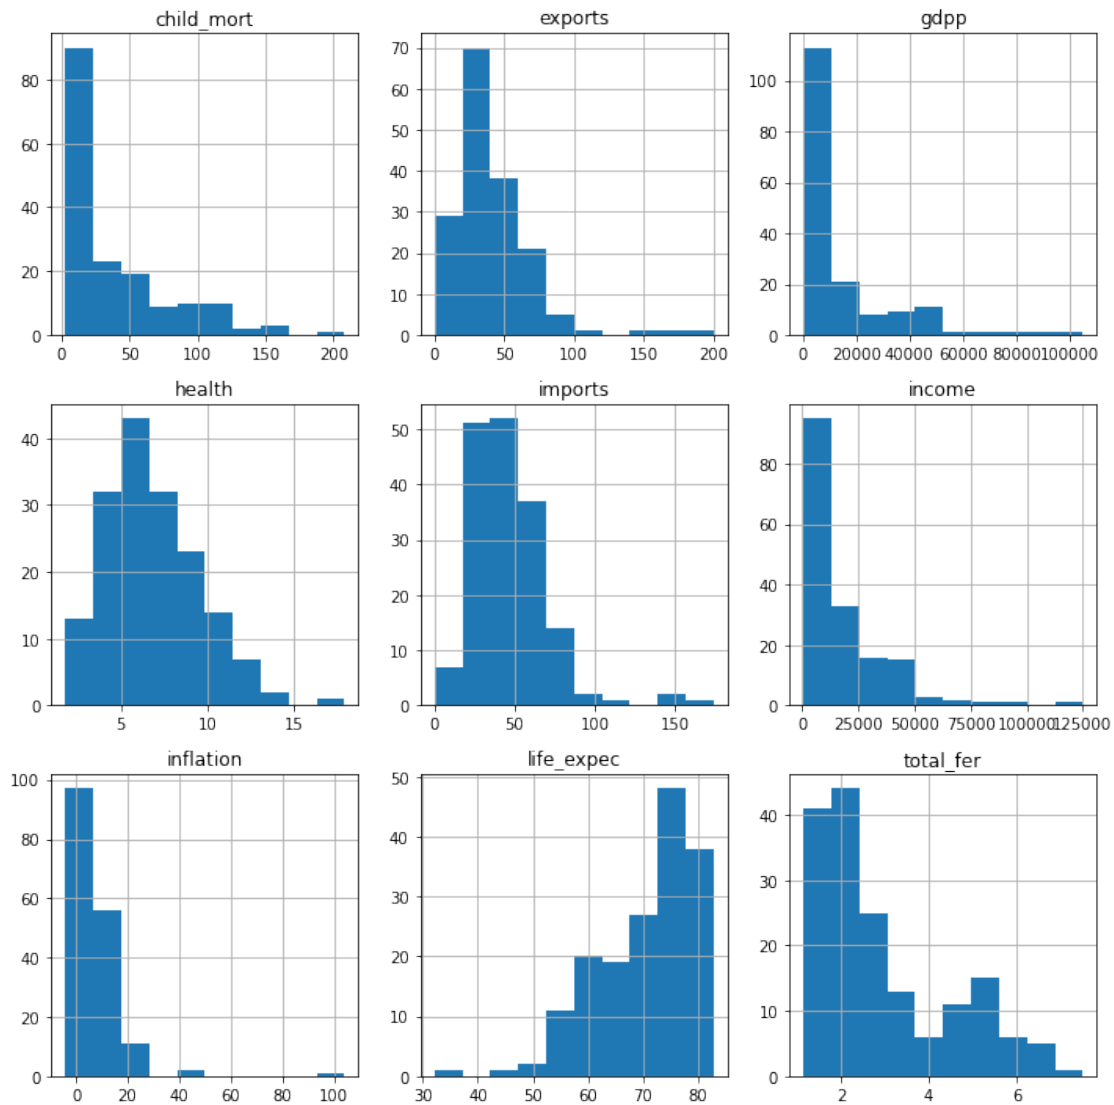
  

	inflation	life_expec	total_fer	gdpp
count	167.000000	167.000000	167.000000	167.000000
mean	7.781832	70.555689	2.947964	12964.155689
std	10.570704	8.893172	1.513848	18328.704809
min	-4.210000	32.100000	1.150000	231.000000
25%	1.810000	65.300000	1.795000	1330.000000
50%	5.390000	73.100000	2.410000	4660.000000
75%	10.750000	76.800000	3.880000	14050.000000
max	104.000000	82.800000	7.490000	105000.000000

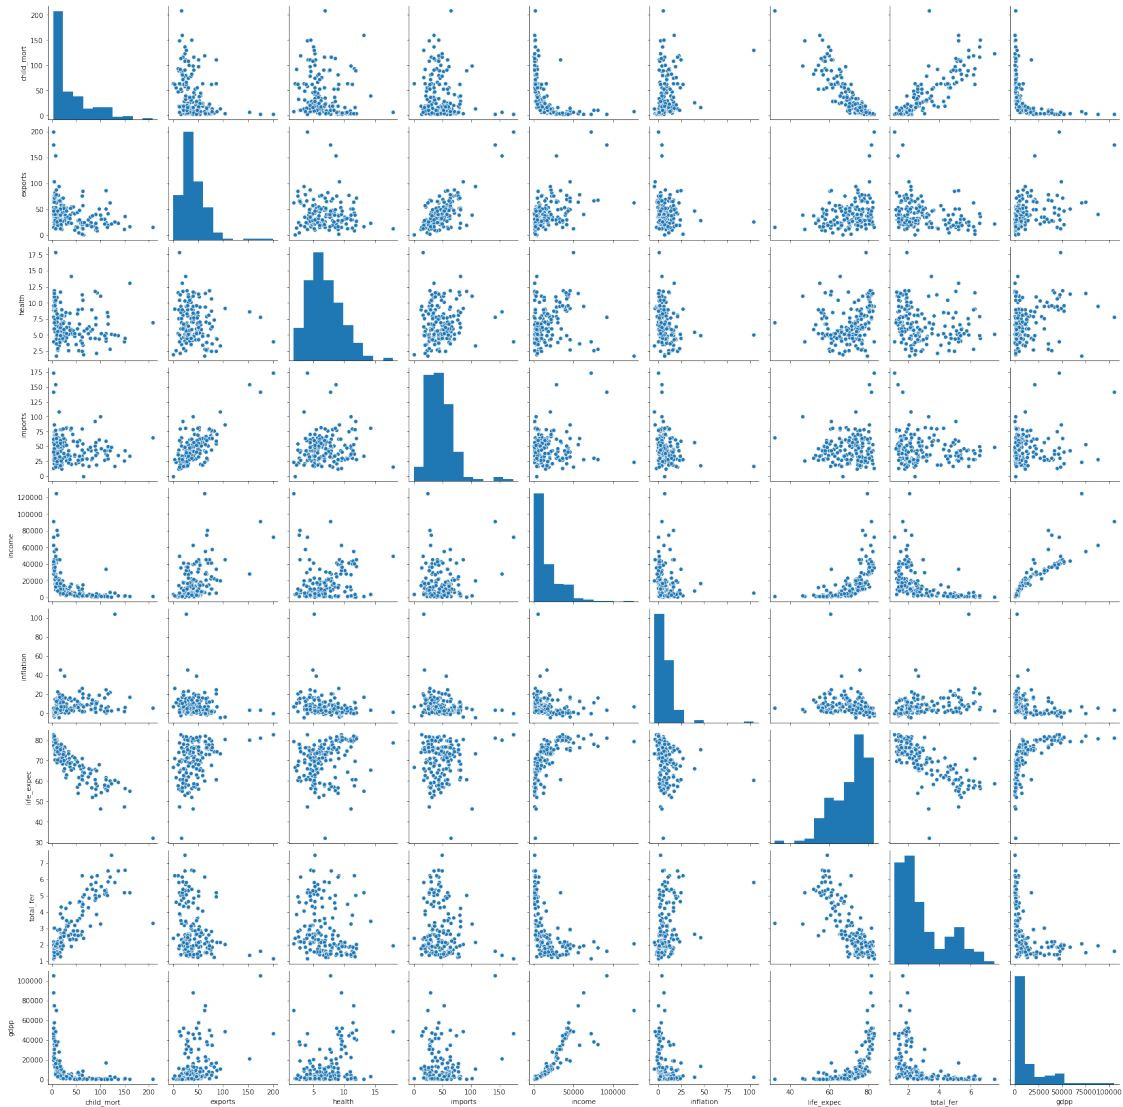
```
[5]: # Checking for missing values
print(CountryData.isnull().sum())
```

```
country      0
child_mort   0
exports      0
health       0
imports      0
income       0
inflation    0
life_expec   0
total_fer    0
gdpp         0
dtype: int64
```

```
[6]: # Histograms to understand distributions
CountryData.hist(figsize=(10,10))
plt.tight_layout()
plt.show()
```



```
[7]: # Pairplot to visualize correlations between variables
sns.pairplot(CountryData.drop("country", axis=1)) # we drop 'country' because
→ it is a non-numeric column
plt.show()
```



```
[8]: # Standardize the data (excluding 'country')
scaler = StandardScaler()
CountryData_scaled = scaler.fit_transform(CountryData.drop('country', axis=1))

# Apply PCA
pca = PCA()
CountryData_pca = pca.fit_transform(CountryData_scaled)

# Print the explained variance ratio
display("Explained variance ratio: ", pca.explained_variance_ratio_)

# Cumulative explained variance
cumulative_variance = np.cumsum(pca.explained_variance_ratio_)
```

```

display("Cumulative explained variance: ", cumulative_variance)

# Plot the explained variance
plt.figure(figsize=(6,4))
plt.bar(range(len(pca.explained_variance_ratio_)), pca.
        ↪explained_variance_ratio_, align='center', label='Individual explained_
        ↪variance')
plt.step(range(len(cumulative_variance)), cumulative_variance,
        ↪where='mid',label='Cumulative explained variance')
plt.ylabel('Explained variance ratio')
plt.xlabel('Principal components')
plt.legend(loc='best')
plt.tight_layout()
plt.show()

```

'Explained variance ratio: '

```

array([0.4595174 , 0.17181626, 0.13004259, 0.11053162, 0.07340211,
       0.02484235, 0.0126043 , 0.00981282, 0.00743056])


```

'Cumulative explained variance: '

```

array([0.4595174 , 0.63133365, 0.76137624, 0.87190786, 0.94530998,
       0.97015232, 0.98275663, 0.99256944, 1.          ])

```



We can see here that if we include the first four components we're explaining approximately 87% of the total variance with only these four components. 87% should be sufficient here so let's go with that to keep our model lean and simple.

### 1.1.2 Select, Build, and Train Model


Let's try out a K-means clustering model because it is simple, flexible, and should work well for only four principal components selected in the previous step.

```
[9]: # Select the first four principal components
CountryData_pca_4 = CountryData_pca[:, :4]

# Determine the number of clusters
# A common technique is the elbow method, where the sum of squared distances to
    ↳ the nearest cluster center
# (within-cluster sum of squares, or WCSS) is calculated for different numbers
    ↳ of clusters, and the number of
# clusters where adding another cluster doesn't significantly improve WCSS is
    ↳ selected.

wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
    kmeans.fit(CountryData_pca_4)
    wcss.append(kmeans.inertia_)
plt.plot(range(1, 11), wcss)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()

# From the plot, we choose the number of clusters where the decrease in WCSS
    ↳ starts to level off (the "elbow")
```



```
[10]: # Suppose the optimal number of clusters is 3
kmeans = KMeans(n_clusters=3, init='k-means++', random_state=42)
clusters = kmeans.fit_predict(CountryData_pca_4)

# Add the cluster assignments to the original dataframe
print(CountryData.shape)
CountryData['Cluster'] = clusters
print(CountryData.shape)

# Calculate mean values for each cluster
cluster_means = CountryData.groupby('Cluster').mean()
display(cluster_means)
```

```
(167, 10)
```

```
(167, 11)
```

	child_mort	exports	health	imports	income	inflation	\
Cluster							
0	91.610417	29.571042	6.433542	43.133333	3897.354167	11.911146	
1	4.897143	58.431429	8.917429	51.508571	45802.857143	2.535000	
2	21.695238	40.484393	6.158333	47.112689	12773.690476	7.608405	

	life_expec	total_fer	gdpp
Cluster			



0	59.239583	4.992083	1909.208333
1	80.245714	1.741143	43117.142857
2	72.984524	2.282738	6717.523810

The mean values above help us interpret the clusters. For example, a cluster with high child mortality and low income and GDP could be considered as a group of countries in severe poverty and in need of aid. In contrast, a cluster with low child mortality and high income and GDP could be considered as a group of wealthy, developed countries.

```
[11]: # Get countries belonging to cluster 0
poverty_cluster = CountryData[CountryData['Cluster'] == 0]

# Sort countries in the poverty cluster by child mortality, in descending order
poverty_cluster_sorted = poverty_cluster.sort_values(by='child_mort',
↳ascending=False)

# Display the sorted countries
print("Number of countries in the poverty cluster: ", poverty_cluster_sorted.
↳shape[0], "\n")
print(poverty_cluster_sorted['country'])
```

Number of countries in the poverty cluster: 48

66	Haiti
132	Sierra Leone
32	Chad
31	Central African Republic
97	Mali
113	Nigeria
112	Niger
3	Angola
25	Burkina Faso
37	Congo, Dem. Rep.
64	Guinea-Bissau
40	Cote d'Ivoire
17	Benin
49	Equatorial Guinea
63	Guinea
28	Cameroon
106	Mozambique
87	Lesotho
99	Mauritania
26	Burundi
116	Pakistan
94	Malawi
150	Togo
0	Afghanistan

```

88          Liberia
36          Comoros
166         Zambia
155         Uganda
56          Gambia
84           Lao
142          Sudan
59          Ghana
147         Tanzania
129         Senegal
38         Congo, Rep.
55          Gabon
126         Rwanda
81         Kiribati
149         Timor-Leste
80          Kenya
93         Madagascar
165         Yemen
108         Namibia
50          Eritrea
137         South Africa
21          Botswana
72           Iraq
136         Solomon Islands
Name: country, dtype: object

```

```
[12]: display(cluster_means)
```

	child_mort	exports	health	imports	income	inflation \
Cluster						
0	91.610417	29.571042	6.433542	43.133333	3897.354167	11.911146
1	4.897143	58.431429	8.917429	51.508571	45802.857143	2.535000
2	21.695238	40.484393	6.158333	47.112689	12773.690476	7.608405

	life_expec	total_fer	gdpp
Cluster			
0	59.239583	4.992083	1909.208333
1	80.245714	1.741143	43117.142857
2	72.984524	2.282738	6717.523810

```

[13]: # Select only the countries from the "needs aid" cluster in the PCA transformed
      ↪ space
aid_cluster_countries_pca = CountryData_pca_4[CountryData['Cluster'] == 0]

# Calculate the Euclidean distance from each country to the centroid of the
      ↪ "needs aid" cluster
aid_cluster_center = kmeans.cluster_centers_[0]

```

```

distances = np.sqrt(((aid_cluster_countries_pca - aid_cluster_center)**2).
    ↳sum(axis=1))

# Convert distances to a pandas series
distances_series = pd.Series(distances,
    ↳index=CountryData[CountryData['Cluster'] == 0].index)

# Select the 10 countries with the smallest distances to the centroid
closest_countries = distances_series.nsmallest(10).index

# Display these countries with their corresponding original data
closest_countries_data = CountryData.loc[closest_countries]
display(closest_countries_data)

```

	country	child_mort	exports	health	imports	income	inflation	\
17	Benin	111.0	23.8	4.10	37.2	1820	0.885	
56	Gambia	80.3	23.8	5.69	42.7	1660	4.300	
166	Zambia	83.1	37.0	5.89	30.9	3280	14.000	
28	Cameroon	108.0	22.2	5.13	27.0	2660	1.910	
94	Malawi	90.5	22.8	6.59	34.9	1030	12.100	
147	Tanzania	71.9	18.7	6.01	29.1	2090	9.250	
106	Mozambique	101.0	31.5	5.21	46.2	918	7.640	
40	Cote d'Ivoire	111.0	50.6	5.30	43.3	2690	5.390	
63	Guinea	109.0	30.3	4.93	43.2	1190	16.100	
59	Ghana	74.7	29.5	5.22	45.9	3060	16.600	

	life_expec	total_fer	gdpp	Cluster
17	61.8	5.36	758	0
56	65.5	5.71	562	0
166	52.0	5.40	1460	0
28	57.3	5.11	1310	0
94	53.1	5.31	459	0
147	59.3	5.43	702	0
106	54.5	5.56	419	0
40	56.3	5.27	1220	0
63	58.0	5.34	648	0
59	62.2	4.27	1310	0

### 1.1.3 Build Model for future datasets

Great! This model seems to be appropriate for the situation. Now lets better construct this model to be used again in the future as these facts about these countries may change from time to time. This is especially true if the countries in most need of aid are the countries always getting the aid.

We basically need to repeat most everything we just did, but this time we need it automated. We need an algorithm that automatically identifies the cluster that corresponds with impoverished

countries. Based on the very strong correlation with impoverished countries and low income per person, I believe a very robust strategy would be to simply select the cluster with the lowest mean income per person.

```
[14]: def cluster_and_identify_countries(df):  
    # Standardize the data (excluding 'country')  
    scaler = StandardScaler()  
    df_normalized = scaler.fit_transform(df.drop('country', axis=1))  
  
    # Apply PCA  
    pca = PCA()  
    df_pca = pca.fit_transform(df_normalized)  
  
    # Select the first four principal components  
    df_pca_4 = df_pca[:, :4]  
  
    # Already determined that the optimal number of clusters is 3  
    kmeans = KMeans(n_clusters=3, init='k-means++', random_state=42)  
    clusters = kmeans.fit_predict(df_pca_4)  
  
    # Add the cluster assignments to the original dataframe  
    df['Cluster'] = clusters  
  
    # Calculate mean values for each cluster  
    cluster_means = df.groupby('Cluster').mean()  
    display(cluster_means)  
  
    # Identify the "needs aid" cluster (the one with the lowest mean income)  
    aid_cluster_id = df.groupby('Cluster')['income'].mean().idxmin()  
    print("The cluster that corresponds to impoverished countries is cluster",  
↪ aid_cluster_id)  
  
    # Filter the dataframe to include only countries in the "needs aid" cluster  
    aid_cluster_countries = df[df['Cluster'] == aid_cluster_id]  
  
    # Calculate the Euclidean distance from each country to the centroid of the  
↪ "needs aid" cluster  
    aid_cluster_center = kmeans.cluster_centers_[aid_cluster_id]  
    distances = np.sqrt(((df_pca_4[df['Cluster'] == aid_cluster_id] -  
↪ aid_cluster_center)**2).sum(axis=1))  
  
    # Add the distances to the dataframe  
    aid_cluster_countries = aid_cluster_countries.assign(Distance=distances)  
  
    # Select the 10 countries with the smallest distances to the centroid  
    top_countries = aid_cluster_countries.nsmallest(10, 'Distance')
```

```
return top_countries
```

```
[15]: # Test the new automated model with the same dataset (we expect the same
      ↪ results)
new_data = pd.read_csv(url_CountryData)
top_countries = cluster_and_identify_countries(new_data)
display(top_countries)
```

	child_mort	exports	health	imports	income	inflation	\
Cluster							
0	91.610417	29.571042	6.433542	43.133333	3897.354167	11.911146	
1	4.897143	58.431429	8.917429	51.508571	45802.857143	2.535000	
2	21.695238	40.484393	6.158333	47.112689	12773.690476	7.608405	

	life_expec	total_fer	gdpp
Cluster			
0	59.239583	4.992083	1909.208333
1	80.245714	1.741143	43117.142857
2	72.984524	2.282738	6717.523810

The cluster that corresponds to impoverished countries is cluster 0

	country	child_mort	exports	health	imports	income	inflation	\
17	Benin	111.0	23.8	4.10	37.2	1820	0.885	
56	Gambia	80.3	23.8	5.69	42.7	1660	4.300	
166	Zambia	83.1	37.0	5.89	30.9	3280	14.000	
28	Cameroon	108.0	22.2	5.13	27.0	2660	1.910	
94	Malawi	90.5	22.8	6.59	34.9	1030	12.100	
147	Tanzania	71.9	18.7	6.01	29.1	2090	9.250	
106	Mozambique	101.0	31.5	5.21	46.2	918	7.640	
40	Cote d'Ivoire	111.0	50.6	5.30	43.3	2690	5.390	
63	Guinea	109.0	30.3	4.93	43.2	1190	16.100	
59	Ghana	74.7	29.5	5.22	45.9	3060	16.600	

	life_expec	total_fer	gdpp	Cluster	Distance
17	61.8	5.36	758	0	0.501046
56	65.5	5.71	562	0	0.549240
166	52.0	5.40	1460	0	0.595864
28	57.3	5.11	1310	0	0.597193
94	53.1	5.31	459	0	0.643892
147	59.3	5.43	702	0	0.697322
106	54.5	5.56	419	0	0.711408
40	56.3	5.27	1220	0	0.806861
63	58.0	5.34	648	0	0.820855
59	62.2	4.27	1310	0	0.869650

## 1.2 Conclusion

In this problem, we utilized Principal Component Analysis (PCA) and K-means clustering to identify countries most in need of aid based on various socio-economic indicators. PCA was instrumental in reducing the dimensionality of our dataset while preserving its essential structure and variations. This allowed us to simplify the dataset, reducing computational complexity, and removing redundant information. The K-means algorithm helped us cluster the countries into different groups based on their socio-economic characteristics. It was then straightforward to identify the cluster of countries with the lowest mean income as the most impoverished and, therefore, the most in need of aid. The distance calculation further provided a means to rank countries within this impoverished cluster according to their closeness to the cluster centroid, i.e., their representative socio-economic condition. The resulting model provided an automated, data-driven approach to aid prioritization, demonstrating the power and utility of unsupervised learning techniques in informing policy decisions. From this analysis, countries like Benin, Gambia, and Zambia emerged as top candidates for aid, highlighting the disparities in global socio-economic conditions and the need for targeted interventions.